

Lecture 2: k-Nearest Neighbors

ECE 2410 – Introduction to Machine Learning

Farzad Farnoud

University of Virginia

Spring 2026

Today's Agenda

- 1 Representing Data
- 2 Matrices, Vectors, and Operations
- 3 kNN
- 4 Distances and Norms
- 5 Python and NumPy
- 6 Practice
- 7 Summary

Last Time: Predict the Major

The Activity

We predicted a student's major based on the majors of their $k = 3$ "closest" neighbors.

Key insight:

- Relies on the assumption that there is a relationship between inputs and outputs
- Similar inputs → similar outputs
- Had we tried to predict cat vs dog person, the error would have been higher

Today

We'll formalize this intuition in our first ML algorithm: **k-Nearest Neighbors (kNN)**

A Bit of Terminology

- **Training data** — solved examples we learn from
- **Test data** — new examples to evaluate on
- **Validation data** — used to choose between approaches and tune model settings
- **Production data** — real-world data after deployment
- **Features (inputs, x)** — measurable inputs (e.g., where you seat)
- **Label (output, y)** — what we want to predict (e.g., major)
- **Generalization** — ability to perform well on *new* data

Key Insight

We care about performance on **test data**, not just training data!

The Classification Problem

Given:

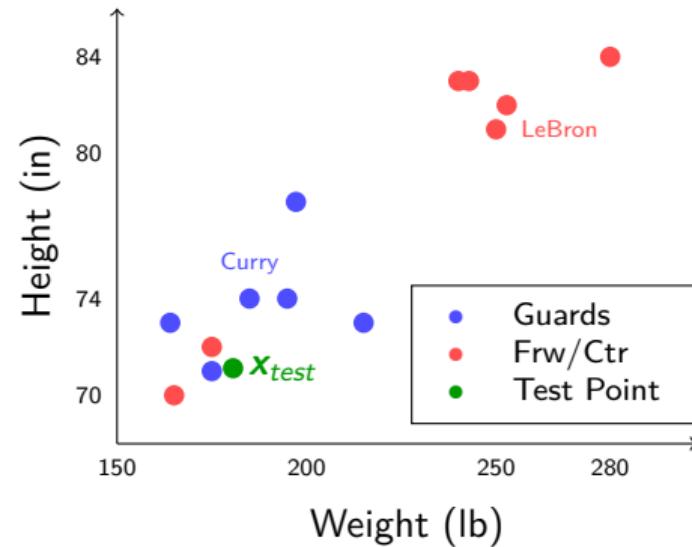
- Training data: n examples
- Each example: (\mathbf{x}_i, y_i)
- \mathbf{x}_i = features (input)
- y_i = label/class (output)

Goal:

- Given a new point \mathbf{x}_{test}
- Predict its label \hat{y}_{test}

Example:

- Predict NBA player position based on height and weight



Is the green point a guard or forward?

Vectors and Matrices: The Basics

Before we proceed: we need a language to describe data and algorithms effectively.

A **Vector** is an ordered list of D numbers
 D is the **dimension**; written as row or column, but it's the same object

$$\mathbf{x} = [x_1 \quad \cdots \quad x_D] \quad \text{or} \quad \mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_D \end{bmatrix}$$

Example: A player's data

$$\mathbf{x} = [250 \quad 81] \quad (\text{Weight, Height})$$

A **matrix** is a rectangular array of numbers with N rows and D columns.

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1D} \\ a_{21} & a_{22} & \cdots & a_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ a_{N1} & a_{N2} & \cdots & a_{ND} \end{bmatrix}$$

We write:

- $\mathbf{x} \in \mathbb{R}^D$
- $\mathbf{A} \in \mathbb{R}^{N \times D}$

From Data to Matrices and Vectors

Raw Data Table:

Name	Weight	Height	Pos
K. Abdul-Jabbar	225	86"	C
M. Abdul-Rauf	162	73"	G
...
L. James	250	81"	F
S. Curry	185	74"	G

Data record i , (x_i, y_i) :

- **Features (x_i):** Measurable properties
($x_{i1} = \text{Height}$, $x_{i2} = \text{Weight}$)
- **Label (y_i):** What we want to predict
(Position)

From Data to Matrices and Vectors

Raw Data Table:

Name	Weight	Height	Pos
K. Abdul-Jabbar	225	86"	C
M. Abdul-Rauf	162	73"	G
...
L. James	250	81"	F
S. Curry	185	74"	G

Data record i , (x_i, y_i) :

- **Features (x_i):** Measurable properties ($x_{i1} = \text{Height}$, $x_{i2} = \text{Weight}$)
- **Label (y_i):** What we want to predict (Position)

Matrices & Vectors:

Row $i \leftrightarrow (x_i, y_i)$

$$\mathbf{X} = \begin{bmatrix} 225 & 86 \\ 162 & 73 \\ \vdots & \vdots \\ 250 & 81 \\ 185 & 74 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} C \\ G \\ \vdots \\ F \\ G \end{bmatrix}$$

Structure:

- $\mathbf{X} \in \mathbb{R}^{N \times D}$ (Features)
- $\mathbf{y} \in \mathbb{R}^N$ (Labels)

The Data Matrix and the Label Vector

Data Matrix: $\mathbf{X} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_N \end{bmatrix} \in \mathbb{R}^{N \times D}$

Label Vector: $\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \in \mathbb{R}^N$

- What are N , D , and row i ?
- y_i = label for \mathbf{x}_i

NBA Example:

$$\mathbf{X} = \begin{bmatrix} 185 & 74 \\ 175 & 72 \\ 250 & 81 \\ \vdots & \vdots \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} G \\ G \\ F \\ \vdots \end{bmatrix}$$

Accessing Elements

Vector Notation:

$$\mathbf{x}_i = [x_{i1} \quad x_{i2} \quad \cdots \quad x_{iD}]$$

- x_{i2} = second element of \mathbf{x}_i =

Matrix Notation:

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

- a_{ij} = element in row i , column j
- Math uses **1-indexing**
- a_{23} = row 2, column 3

Extracting from Matrices:

- Row i of A : $A_{i,:}$ or A_i
- Column j of A : $A_{:,j}$

Example:

$$\mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- $a_{12} = 2$
- Row 1: $A_{1,:} = [1 \quad 2 \quad 3]$
- Column 2: $A_{:,2} = \begin{bmatrix} 2 \\ 5 \end{bmatrix}$

In-Class Activity: Practice the Notation!

Dataset: Student Study Habits

Name	Hours/Week	Coffees/Day	Grade
Alice	15	2	A
Bob	8	1	B
Carol	20	4	A
Dave	5	0	C

Task: Considering the (X, y) representation, for each quantity:

- Describe it in words
- Write it out mathematically

Example: x_1

- Alice's feature vector
- $[15 \quad 2]$

Your turn: Solve the following individually, then discuss with your neighbor. Revise your work and turn in the final answers.

- ① x_2
- ② $x_{3,2}$
- ③ X
- ④ y
- ⑤ y_1
- ⑥ $X_{:,1}$

In-Class Activity: Solutions

Answers:

① $x_2 = \text{Bob's feature vector}$
 $[8 \ 1]$

② $x_{3,2} = \text{Carol's coffee count}$
4

③ $\mathbf{X} = \text{Feature matrix (all data)}$
$$\begin{bmatrix} 15 & 2 \\ 8 & 1 \\ 20 & 4 \\ 5 & 0 \end{bmatrix}$$

④ $\mathbf{y} = \text{Label vector (all grades)}$

$$\begin{bmatrix} \text{A} \\ \text{B} \\ \text{A} \\ \text{C} \end{bmatrix}$$

⑤ $y_1 = \text{Alice's grade}$
A

⑥ $\mathbf{X}_{:,1} = \text{Study hours column}$
$$\begin{bmatrix} 15 \\ 8 \\ 20 \\ 5 \end{bmatrix}$$

The Inner (dot) Product

Definition

The **inner product** (or dot product) of two vectors $\mathbf{a}, \mathbf{b} \in \mathbb{R}^D$ is the sum of element-wise products:

$$\langle \mathbf{a}, \mathbf{b} \rangle = \sum_{j=1}^D a_j b_j$$

Interpretation: Measures **similarity** or alignment. *Higher dot product = more similar!*

Example: Movie Ratings $\begin{bmatrix} \text{Action} \\ \text{Romance} \end{bmatrix}$: $\mathbf{a} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}$, $\mathbf{b} = \begin{bmatrix} 4 \\ 2 \end{bmatrix}$, $\mathbf{c} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$

$$\langle \mathbf{a}, \mathbf{b} \rangle = 5(4) + 1(2) = 22 \text{ High similarity!} \quad \langle \mathbf{a}, \mathbf{c} \rangle = 5(1) + 1(5) = 10 \text{ Low similarity.}$$

Matrix-Vector Multiplication: $\mathbf{A}\mathbf{x}$

Rule for $\mathbf{A}\mathbf{x}$

Multiply column i of \mathbf{A} by x_i and sum.

Example:

$$\begin{bmatrix} 1 & -1 \\ 3 & -4 \\ -1 & 2 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = 2 \cdot \begin{bmatrix} 1 \\ 3 \\ -1 \end{bmatrix} + 1 \cdot \begin{bmatrix} -1 \\ -4 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 0 \end{bmatrix}$$

Matrix-Vector Multiplication: A Recipe

Example: Party Planning

- We want to make **Tacos** and **Pizza**.
- **Matrix A (Recipes):** Ingredients needed per item.
- **Vector x (Menu):** How many of each we want.

$$\begin{matrix} \text{Pizza} & \text{Taco} \\ \text{Cheese} & \text{Ground Beef} \\ \text{Rice} & \text{Tortillas} \\ \text{Salsa} & \end{matrix} \begin{bmatrix} 2 & 1 \\ 1 & 3 \\ 0.5 & 1 \end{bmatrix} \times \begin{matrix} \text{Pizza} & \text{Taco} \\ \text{Cheese} & \text{Ground Beef} \\ \text{Rice} & \text{Tortillas} \\ \text{Salsa} & \end{matrix} \begin{bmatrix} 5 \\ 10 \end{bmatrix} = 5 \cdot \begin{bmatrix} 2 \\ 1 \\ 0.5 \end{bmatrix} + 10 \cdot \begin{bmatrix} 1 \\ 3 \\ 1 \end{bmatrix} = \begin{bmatrix} 10 \\ 5 \\ 2.5 \end{bmatrix} + \begin{bmatrix} 10 \\ 30 \\ 10 \end{bmatrix} = \begin{bmatrix} 20 \\ 35 \\ 12.5 \end{bmatrix} \begin{matrix} \text{Cheese} & \text{Ground Beef} \\ \text{Rice} & \text{Tortillas} \\ \text{Salsa} & \end{matrix}$$

Result: Total shopping list!

Vector-Matrix Multiplication: xA

Rule for xA

Multiply element i of the vector with row i of the matrix and sum.

Example: Government Budget

- **Vector x (Investments):** Budget allocated to each sector.
- **Matrix A (Impact):** Each investment's contribution to outcomes.

$$\begin{bmatrix} \text{books} & \text{hospital} & \text{house} \\ 2 & 3 & 1 \end{bmatrix} \times \begin{bmatrix} \text{smile} & \text{shield} & \text{dollar} \\ 3 & 1 & 2 \\ 2 & 4 & 1 \\ 1 & 2 & 5 \end{bmatrix} = [6 \ 2 \ 4] + [6 \ 12 \ 3] + [1 \ 2 \ 5] = \begin{bmatrix} \text{smile} & \text{shield} & \text{dollar} \\ 13 & 16 & 12 \end{bmatrix}$$

Total impact: 13 happiness, 16 security, 12 wealth units

Recall: Generalization

Generalization

The ability of a model to perform well on **new, unseen data** (test set), not just the data it was trained on.

- **Training Set:** Examples given to the model to learn from (e.g., historical player data).
- **Test Set:** New examples hidden from the model during training, used to evaluate performance (e.g., new rookies).

We care about how well we predict the unseen, not memorize observations!

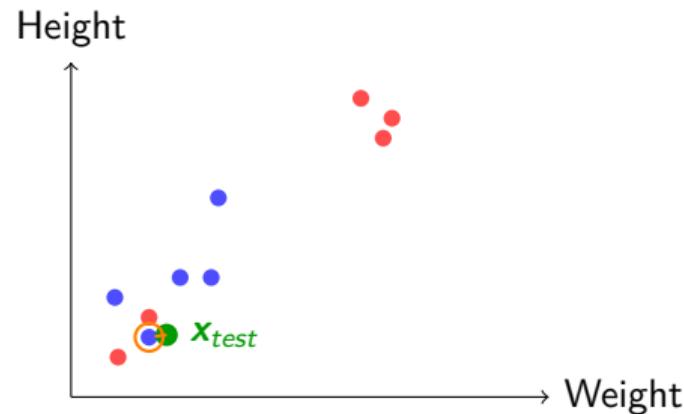
1-Nearest Neighbor (1-NN)

Algorithm

Find the training point closest to x_{test} and copy its label.

Steps:

- ① Compute distance from x_{test} to every training point
- ② Find the training point with minimum distance
- ③ Assign its label to x_{test}



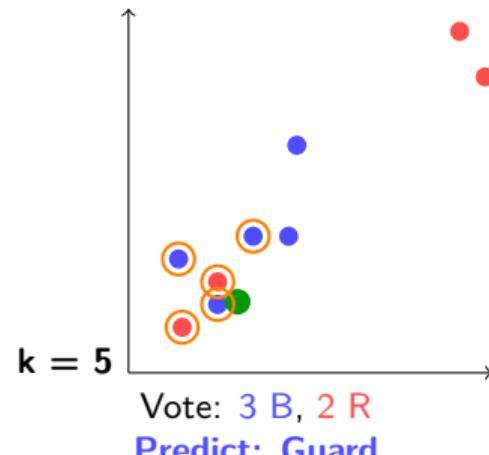
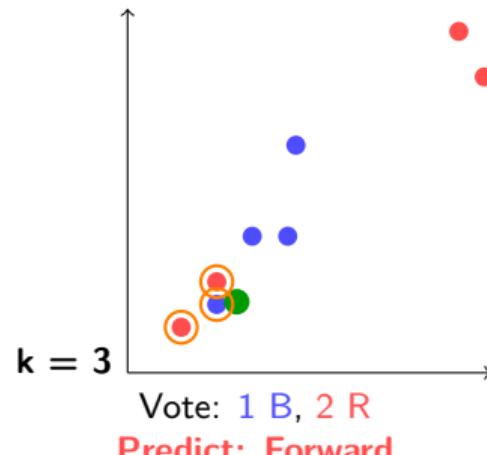
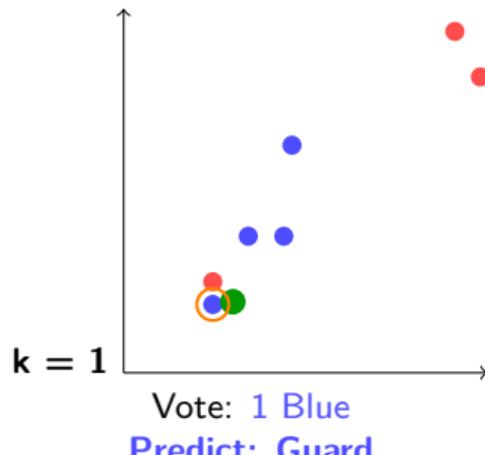
Predict: Guard

k-Nearest Neighbors (k-NN)

Algorithm

Find the k closest training points and take a **majority vote**.

Scenario: Test player at 180 lbs, 71 inches. Guard or Forward?



Later: How do we choose k ?

The Argmin Operation

Definition

$\arg \min_i f(i)$ returns the **index** i that minimizes $f(i)$, not the minimum value itself.

Example: Let distances = [5.2, 3.1, 8.7, 2.4, 6.0]

- $\min(\text{distances}) = 2.4$ (the smallest value)
- $\arg \min(\text{distances}) = 3$ (the index of 2.4, using 0-indexing)

In Python/NumPy:

- `np.min(distances) → 2.4`
- `np.argmin(distances) → 3`

For kNN

We need argmin to find *which* training point is closest, then look up its label.

How Do We Measure “Closest”?

Key Question

What does it mean for two points to be “close”?

Consider two points:

$$\mathbf{x} = (x_1, x_2, \dots, x_d) \quad \text{and} \quad \mathbf{z} = (z_1, z_2, \dots, z_d)$$

Common distance metrics:

Euclidean (L2) Distance:

$$d_2(\mathbf{x}, \mathbf{z}) = \sqrt{\sum_{j=1}^d (x_j - z_j)^2}$$

“Straight line” distance

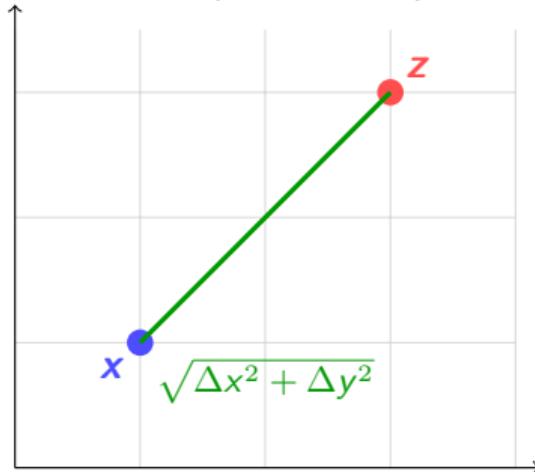
Manhattan (L1) Distance:

$$d_1(\mathbf{x}, \mathbf{z}) = \sum_{j=1}^d |x_j - z_j|$$

“City block” distance

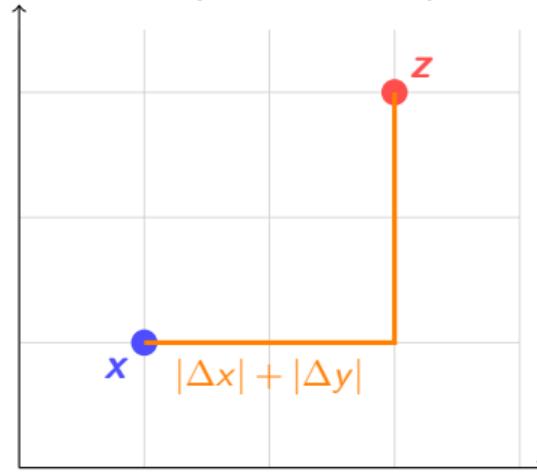
Visualizing L1 vs L2 Distance

L2 (Euclidean)



Straight Line

L1 (Manhattan)



Grid Path

The General L_p Norm

Norm is the distance from the origin. Measures the size of a vector.

$$\|\mathbf{x}\| = d(\mathbf{x}, \mathbf{0})$$

L_p Norm

$$\|\mathbf{x}\|_p = \left(\sum_{j=1}^d |x_j|^p \right)^{1/p}$$

Distance is the norm of the difference vector.

$$d(\mathbf{x}, \mathbf{z}) = \|\mathbf{x} - \mathbf{z}\|$$

L_p Distance

$$d_p(\mathbf{x}, \mathbf{z}) = \left(\sum_{j=1}^d |x_j - z_j|^p \right)^{1/p}$$

We usually talk about norms for simplicity.

Common L_p Norms

L_p Norm

$$\|\mathbf{x}\|_p = \left(\sum_{j=1}^d |x_j|^p \right)^{1/p}$$

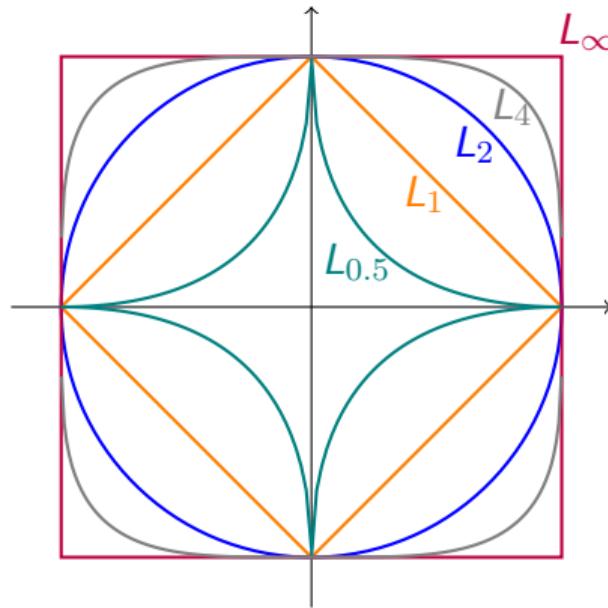
- $p = 1$ (Manhattan / Taxicab): $\|\mathbf{x}\|_1 = |x_1| + |x_2| + \cdots + |x_d|$
- $p = 2$ (Euclidean): $\|\mathbf{x}\|_2 = \sqrt{x_1^2 + x_2^2 + \cdots + x_d^2} = \sqrt{\langle \mathbf{x}, \mathbf{x} \rangle}$ (square root of the inner product of the vector with itself)
- $p = \infty$ (Max / Chebyshev): $\|\mathbf{x}\|_\infty = \max_j |x_j|$

Example: $\mathbf{x} = [3 \quad -4]$

- $\|\mathbf{x}\|_1 = |3| + |-4| = 7$
- $\|\mathbf{x}\|_2 = \sqrt{9 + 16} = 5$
- $\|\mathbf{x}\|_\infty = \max(3, 4) = 4$

Unit Circles (Contours of Norms)

Each curve represents the set of points at a unit distance away from the origin in L_p metric.



L2 (Euclidean):

- **Rotationally Invariant:** Distance doesn't change if you rotate the axes.
- "As the crow flies"

L1 (Manhattan):

- **Axis Dependent:** Axes have specific, distinct meanings (e.g., Age vs. Salary).
- Less susceptible to outliers in one dimension (each axis contributes linearly).

Why NumPy?

Without NumPy:

- Python lists are slow for math
- Must write explicit loops
- Memory inefficient

With NumPy:

- Fast vectorized operations
- Clean, readable code
- Built-in math functions

Example: Add two vectors

Python lists:

```
[a[i]+b[i] for i in range(n)]
```

NumPy arrays:

```
a + b
```

NumPy is often **100x faster!**

Computing Distances in NumPy

The `np.linalg.norm` function:

```
np.linalg.norm(x - z, ord=2)
```

- Computes the L_p norm of the difference vector
- `ord=2` (default): Euclidean distance
- `ord=1`: Manhattan distance
- `ord=np.inf`: Maximum absolute difference

Example:

```
x = np.array([3, 4])
z = np.array([0, 0])
np.linalg.norm(x - z) → 5.0
```

(This is $\sqrt{3^2 + 4^2} = \sqrt{25} = 5$)

Today's Notebook Activities

Prerequisites from HW01:

- §7: NumPy arrays, vectors, matrices
- §7.3: Slicing and indexing ($M[i, :]$)
- §8: Boolean indexing, `np.argmax`

Today's NEW activities (in notebook):

- ① **Activity 1: Computing Distances**
 - Use `np.linalg.norm` to compute Euclidean distance
- ② **Activity 2: Finding the Nearest Neighbor**
 - Compute distances to all training points
 - Use `np.argmin` to find the nearest
- ③ **Activity 3: Classify All Test Points**
 - Implement 1-NN for multiple test points

Note

Complete HW01 §7–8 first if you haven't already!

Key Takeaways

① k-Nearest Neighbors:

- Find k closest training points
- Take majority vote to classify
- Simple but powerful!

② Distance Metrics:

- L2 (Euclidean): straight-line distance
- L1 (Manhattan): city-block distance
- Use `np.linalg.norm(x-z, ord=p)`

③ NumPy Essentials:

- Arrays have `.shape` and `.dtype`
- Indexing: `M[row, col]`, slicing with `:`
- Vectorized operations are fast!

Next Time

We'll implement kNN from scratch and apply it to real data!

k-Nearest Neighbors Algorithm:

- ① Find k closest training points to test point
- ② Take majority vote of their labels
- ③ Prediction depends on distance metric and k

Key Concepts:

- Vectors represent data points
- Matrices organize datasets (rows = samples, columns = features)
- Distance metrics: L1 (Manhattan), L2 (Euclidean), L p
- Matrix operations: Ax (combines columns), xA (combines rows)